

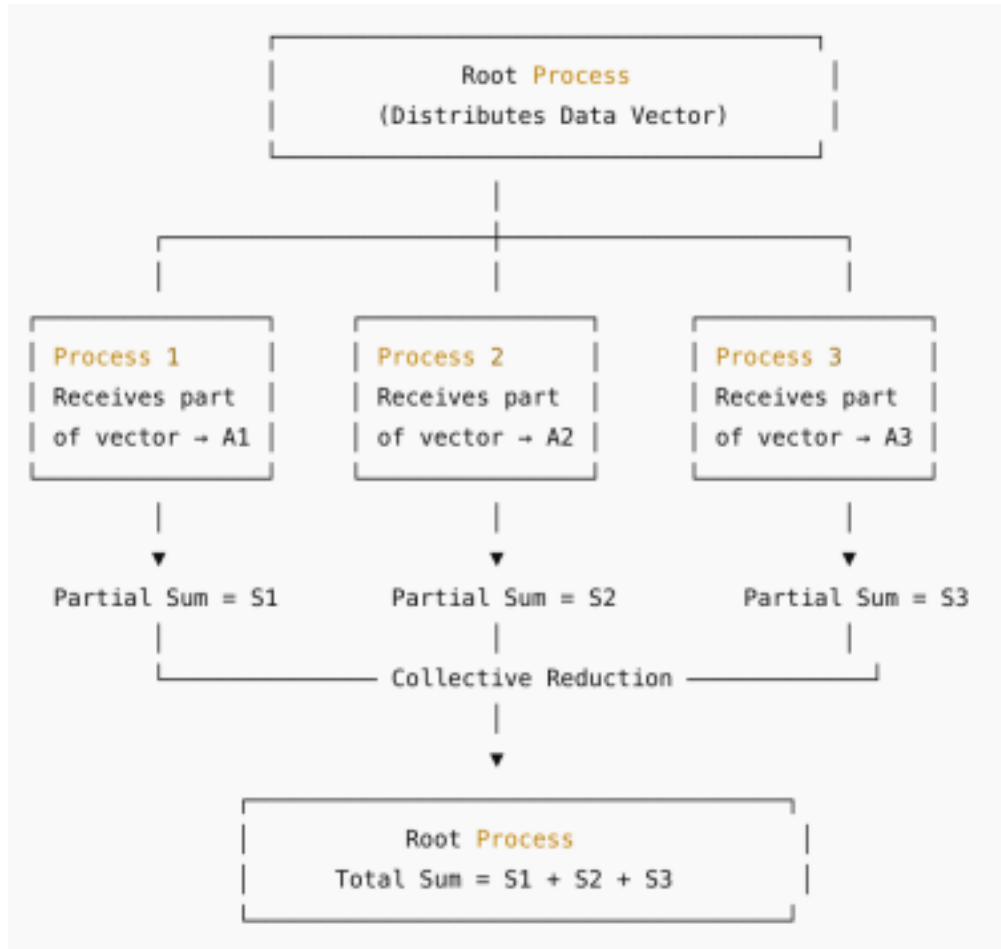
❖❖ 🎲 Lab Assignment #5 – Distributed Partial Summation using MPI

❖❖ 🎯 Objective

This lab focuses on dividing a large computational task (sum of a big vector) across multiple nodes (processes). Each node will handle a subset of data, compute a **partial sum**, and then combine the results at the **root process** using **collective MPI operations**.

❖❖ 🧠 Concepts Covered

- Data decomposition (dividing large datasets)
- Inter-process communication (MPI)
- Reduction and collective operations (`MPI_Reduce`, `MPI_Gather`)
- Understanding root and worker processes



🔍 🌱 Step-by-Step Tasks

Task 1: Vector Initialization

- Create a large 1D vector A of size $N = 10,000,000$ on the **root process**. • Initialize it with values $A[i] = i + 1$.

Task 2: Data Distribution

- Divide A into equal sub-vectors and distribute them among all processes using: • `MPI_Scatter()`
- Each process receives its sub-array `local_A`.

Task 3: Local Computation

- Each process computes the **partial sum** of its sub-vector:
- `double local_sum = 0;`

- `for (int i = 0; i < local_size; i++)`
- `local_sum += local_A[i];`

Task 4: Reduction (Collective Operation)

- Use **MPI_Reduce()** to combine all partial sums into one **total sum** at the root node:


```
MPI_Reduce(&local_sum, &global_sum, 1, MPI_DOUBLE, MPI_SUM, 0,
MPI_COMM_WORLD);
```

Task 5: Display Results

- The **root process** prints the final sum and verifies correctness with the expected result:

$$\text{Expected Sum} = \frac{N(N+1)}{2}$$

$$\text{Expected Sum} = 2N(N+1)$$

Sample Code (C using MPI)

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char* argv[]) {
    int rank, size;
    long N = 100000000;
    double *A = NULL;
    double local_sum = 0.0, global_sum = 0.0;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    long local_n = N / size;
    double *local_A = (double*)malloc(local_n * sizeof(double));

    if (rank == 0) {
        A = (double*)malloc(N * sizeof(double));
        for (long i = 0; i < N; i++)
            A[i] = i + 1;
    }

    // Distribute subarrays to each process
    MPI_Scatter(A, local_n, MPI_DOUBLE, local_A, local_n, MPI_DOUBLE, 0,
MPI_COMM_WORLD);

    // Local computation
    for (long i = 0; i < local_n; i++)
        local_sum += local_A[i];
```

```

// Reduction
MPI_Reduce(&local_sum, &global_sum, 1, MPI_DOUBLE, MPI_SUM, 0,
MPI_COMM_WORLD);

if (rank == 0) {
double expected = (N * (N + 1)) / 2.0;
printf("Total Sum = %.0f | Expected = %.0f | Difference = %.5f\n",
global_sum, expected, expected - global_sum); free(A);
}

free(local_A);
MPI_Finalize();
return 0;
}

```

Discussion Questions

1. What happens if the vector size **N** is not divisible by the number of processes?
2. How can you modify the program to handle **uneven partitions**?
3. How would performance differ between using `MPI_Reduce` vs. `MPI_Gather` + local summation?
4. How could this same approach be extended to **matrix summation or averaging**?

Bonus Challenge

Modify the program to:

- Compute both **sum and average** using a single collective operation (`MPI_Allreduce`).
- Measure execution time using `MPI_Wtime()` and compare with serial CPU computation.

Submission: Code submission will be accepted only on GitHub Public repo